


# Understanding LLM from Methods to Practices

---

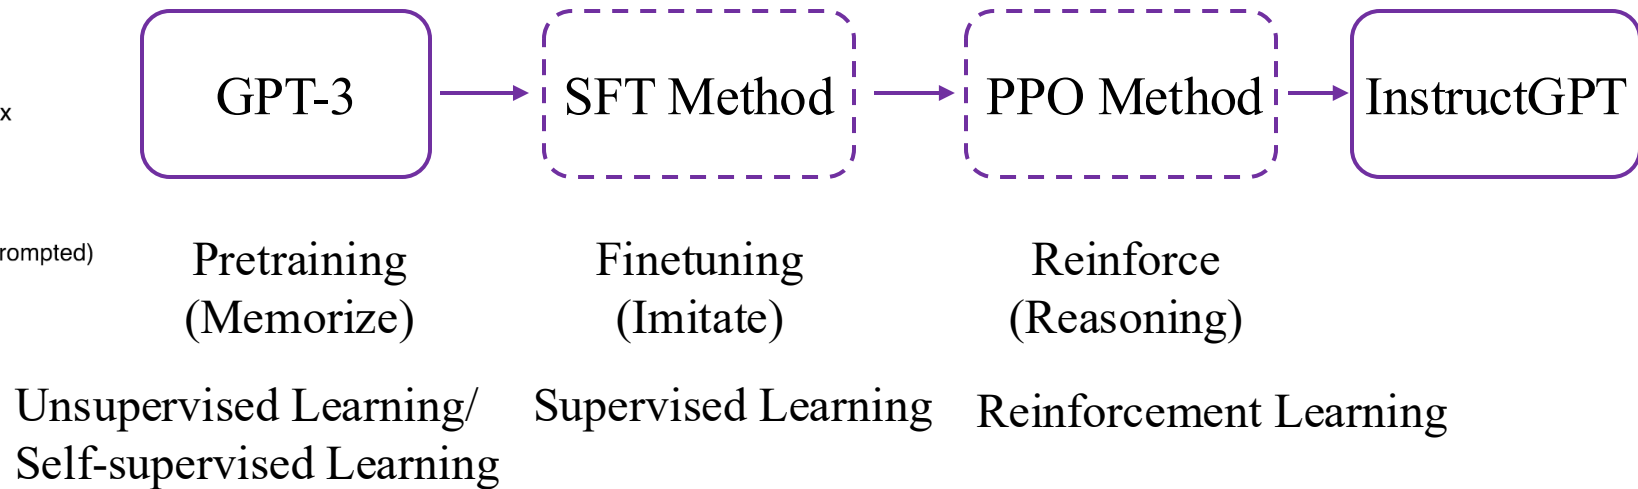

# Scaling Effectiveness of LLM



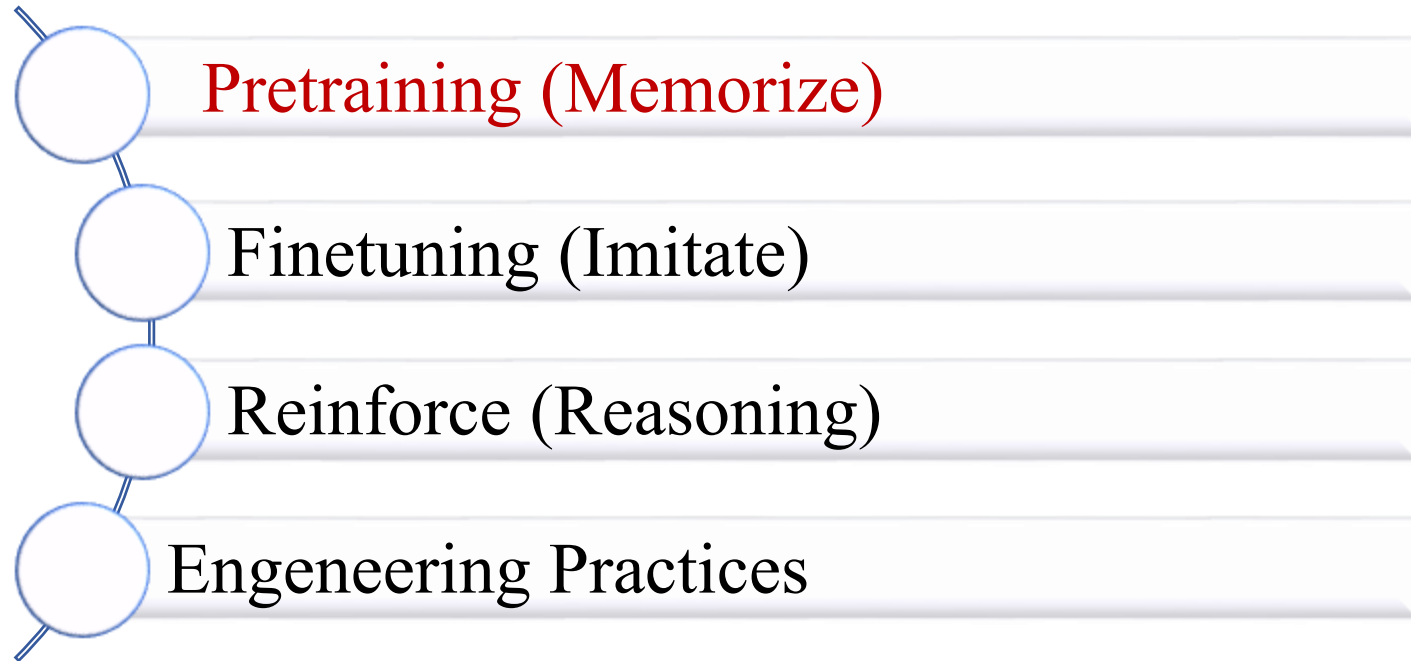
Three observations

- Scaling model size contributes to the performance
- Supervised FineTuning (SFT) and Proximal Policy Optimization (PPO) contribute to the performance
- 1.3B GPT-3 with SFT surpasses the vanilla 175B GPT-3

# Scaling Effectiveness of LLM



# Outline


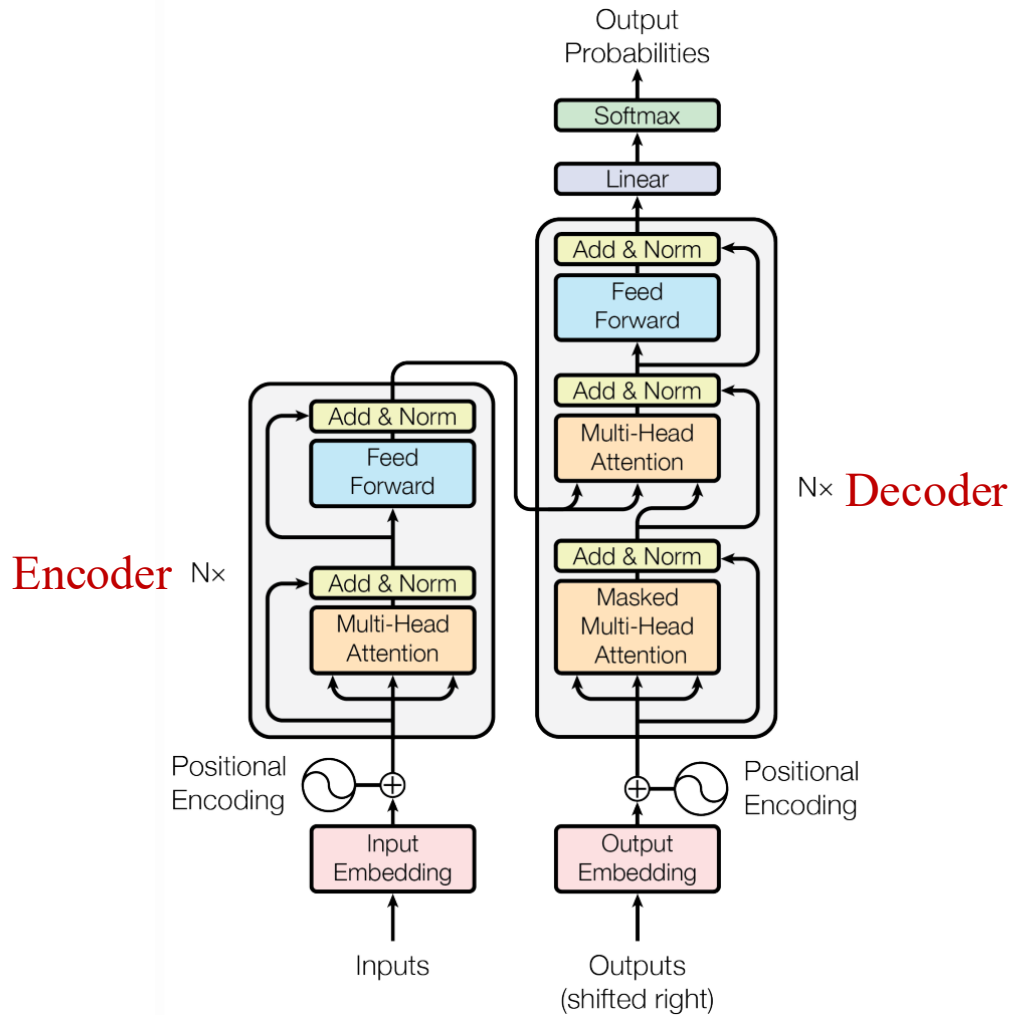


# Encoder-decoder Architecture of Transformer

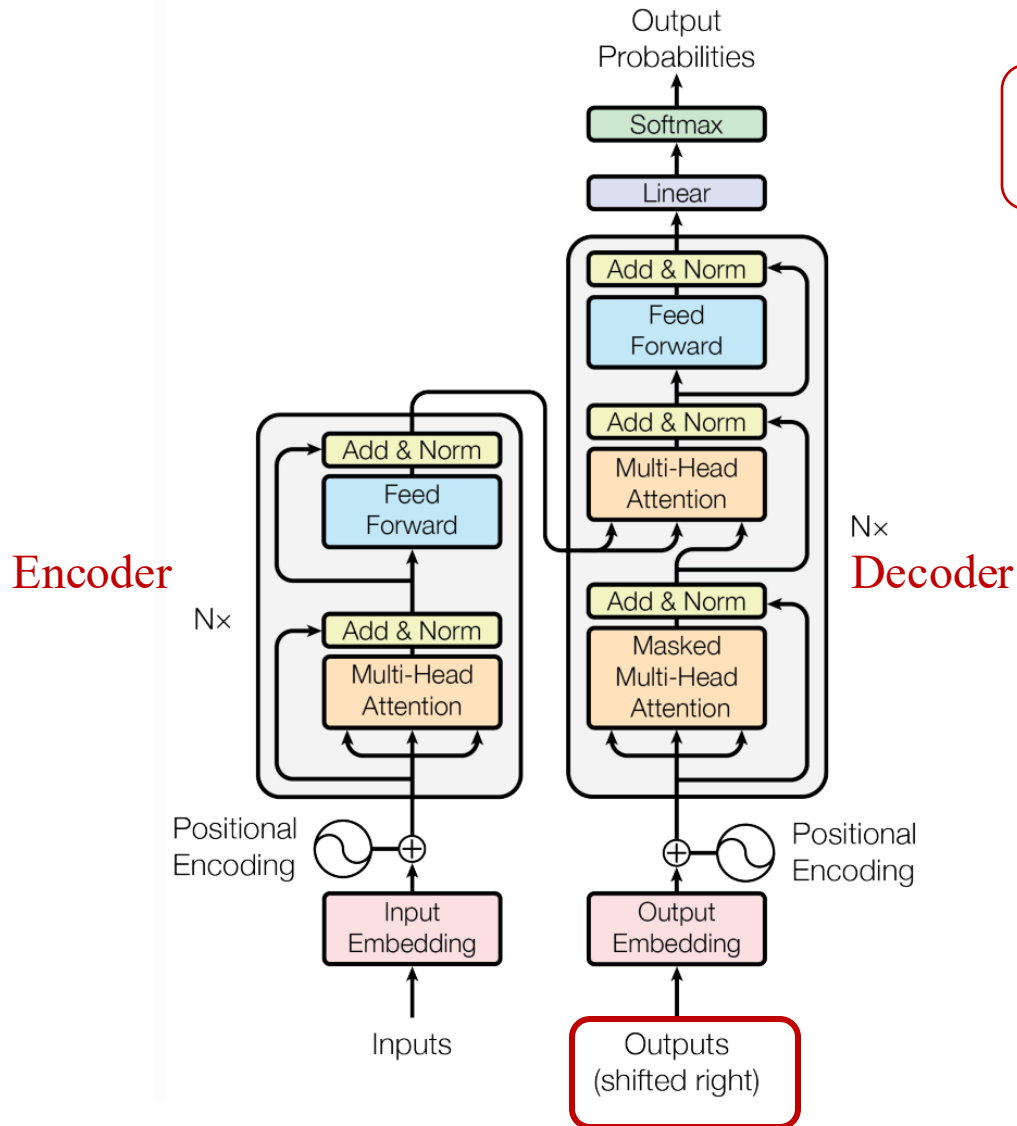
Goal of **encoder**: get contextualized representation for each token

Self-attention in **encoder**

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



# Encoder-decoder Architecture of Transformer



Goal of **decoder**: accomplish the next token prediction task autoregressively

Suppose the target sequence is ["A", "B", "C"]

**input** of decoder: [<SOS>, "A", "B", "C"]

**label**: ["A", "B", "C", <EOS>]

Task of decoder:

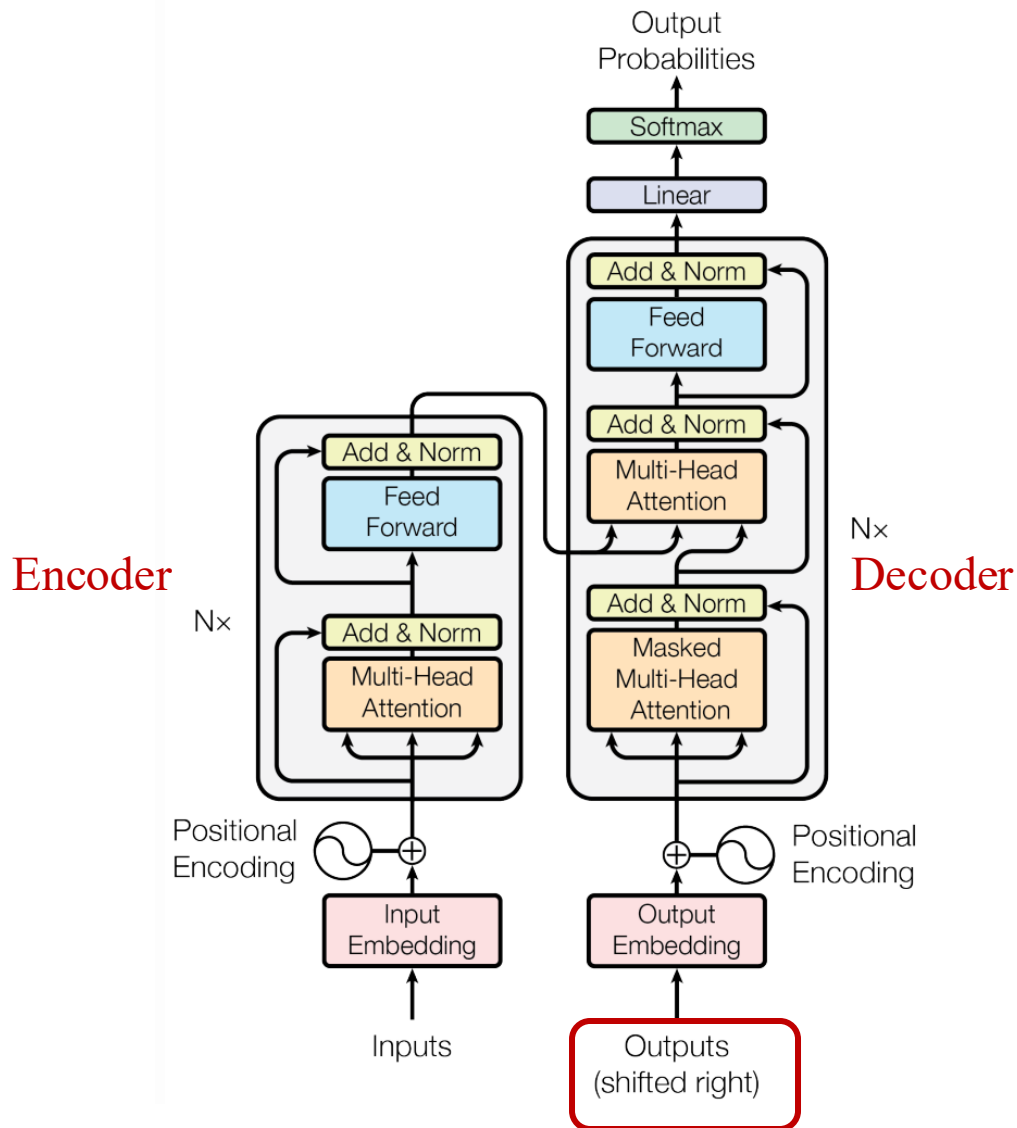
input [<SOS>], predict "A"

input [<SOS>, "A"], predict "B"

input [<SOS>, "A", "B"], predict "C"

input [<SOS>, "A", "B", "C"], predict <EOS>

# Encoder-decoder Architecture of Transformer



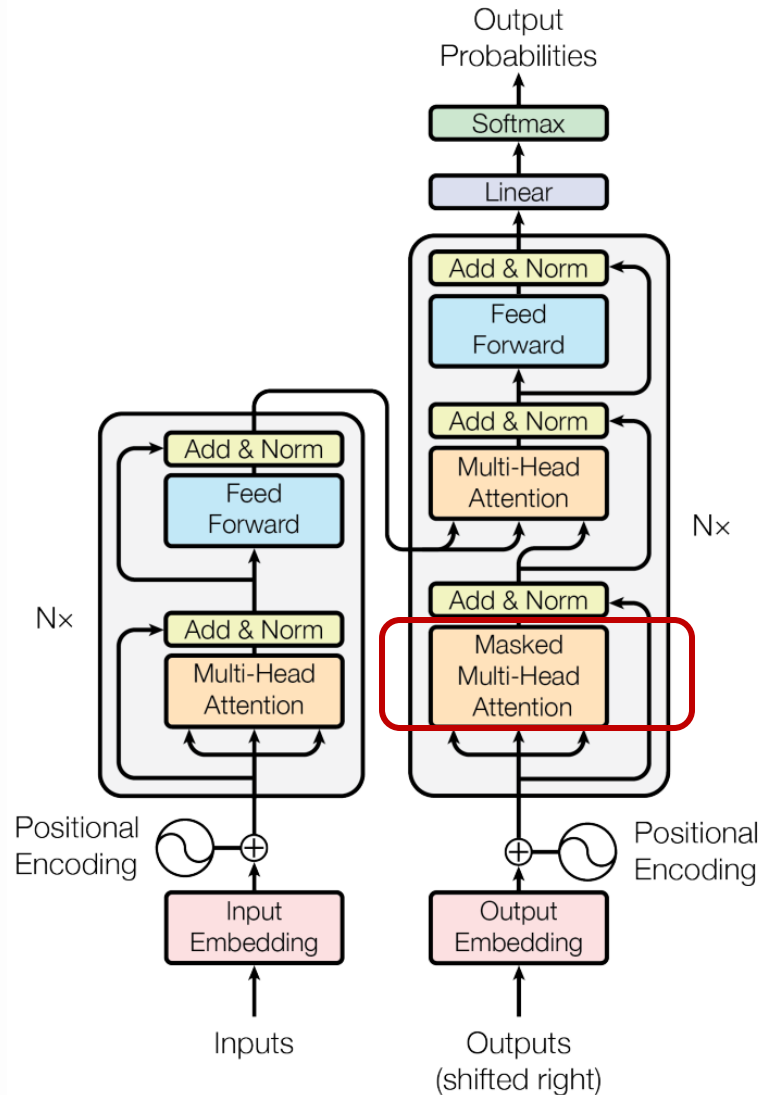
Next token prediction task in **decoder** part

**input** of decoder:  $[\langle \text{SOS} \rangle, x_1, x_2, \dots, x_T]$       SOS: start of sentence


**label:**  $[x_1, x_2, x_3, \dots, x_T, \langle \text{EOS} \rangle]$       EOS: end of sentence

- Vanilla implementation: Given previous  $t - 1$  tokens, predict the  $t$ -th token
- Efficient implementation: Given all  $T$  tokens as inputs, predict all the outputs simultaneously using **masked self-attention**
  - Benefits: Fully exploits **parallel** computation capabilities

# Encoder-decoder Architecture of Transformer



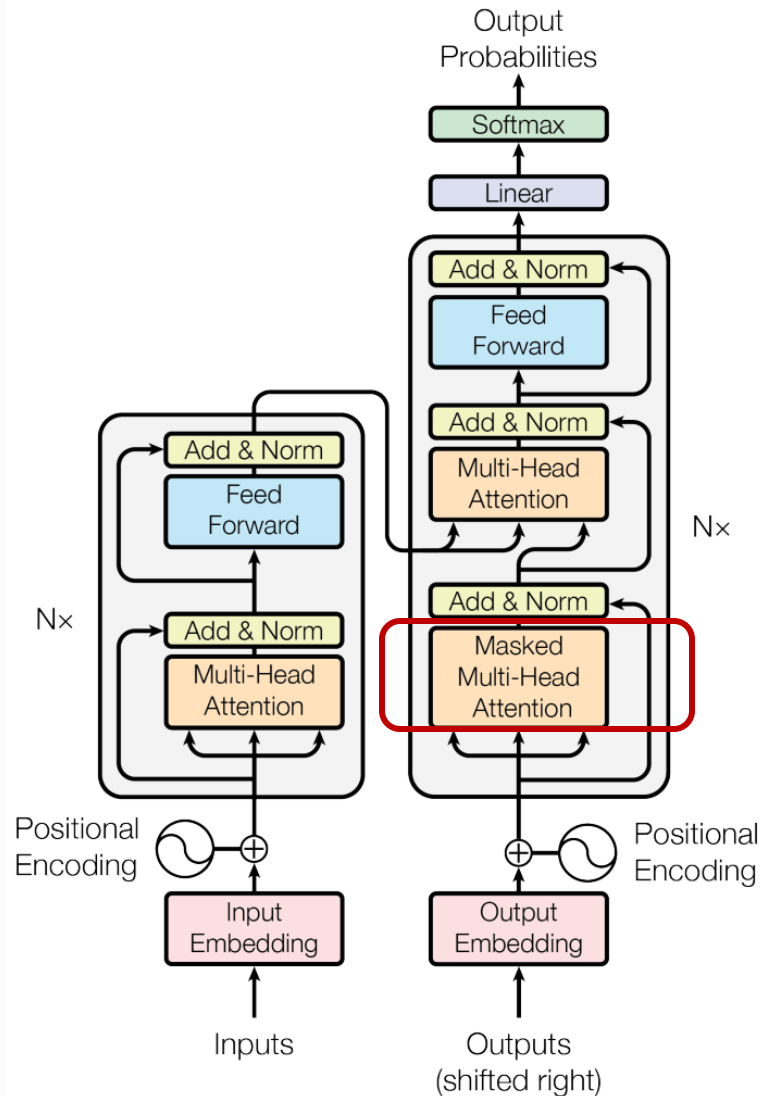
Masked self-attention in decoder



■  $h_t$  is only related to  $[x_1, x_2, \dots, x_t]$



# Encoder-decoder Architecture of Transformer



Efficient implementation of **masked** self-attention in **decoder**

$$\text{MaskedAttention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}} + \mathbf{M}\right)V$$

- $\mathbf{M}$  is **causal mask** with upper triangular part is  $-\infty$ , lower triangular part is 0, for example:

$$\mathbf{M} = \begin{bmatrix} 0 & -\infty & -\infty & -\infty \\ 0 & 0 & -\infty & -\infty \\ 0 & 0 & 0 & -\infty \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

# Encoder-decoder Architecture of Transformer

$$\text{MaskedAttention}(Q, K, V) = \text{softmax}\left(\underbrace{\frac{QK^T}{\sqrt{d_k}}}_S + M\right)V \quad M = \begin{bmatrix} 0 & -\infty & -\infty & -\infty \\ 0 & 0 & -\infty & -\infty \\ 0 & 0 & 0 & -\infty \\ 0 & 0 & 0 & 0 \end{bmatrix} \quad \text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}} \quad e^{-\infty} = 0$$

$$S = \frac{QK^T}{\sqrt{d_k}} = \begin{bmatrix} s_{11} & s_{12} & s_{13} & s_{14} \\ s_{21} & s_{22} & s_{23} & s_{24} \\ s_{31} & s_{32} & s_{33} & s_{34} \\ s_{41} & s_{42} & s_{43} & s_{44} \end{bmatrix} \quad S + M = \begin{bmatrix} s_{11} & -\infty & -\infty & -\infty \\ s_{21} & s_{22} & -\infty & -\infty \\ s_{31} & s_{32} & s_{33} & -\infty \\ s_{41} & s_{42} & s_{43} & s_{44} \end{bmatrix} \quad \text{softmax}(S + M) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ c_{21} & c_{22} & 0 & 0 \\ c_{31} & c_{32} & c_{33} & 0 \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$

$$\text{softmax}(S + M)V = \begin{bmatrix} 1 & 0 & 0 & 0 \\ c_{21} & c_{22} & 0 & 0 \\ c_{31} & c_{32} & c_{33} & 0 \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \end{bmatrix} = \begin{bmatrix} v_1 \\ c_{21}v_1 + c_{22}v_2 \\ c_{31}v_1 + c_{32}v_2 + c_{33}v_3 \\ c_{41}v_1 + c_{42}v_2 + c_{43}v_3 + c_{44}v_4 \end{bmatrix} = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix}$$

- $v_t$  is row vector
- $h_t$  is only related to tokens before  $t$

# Encoder-decoder Architecture of Transformer

Goal of **cross attention**: compute the relationship between  $h_t$  from decoder and contextualized representation from encoder

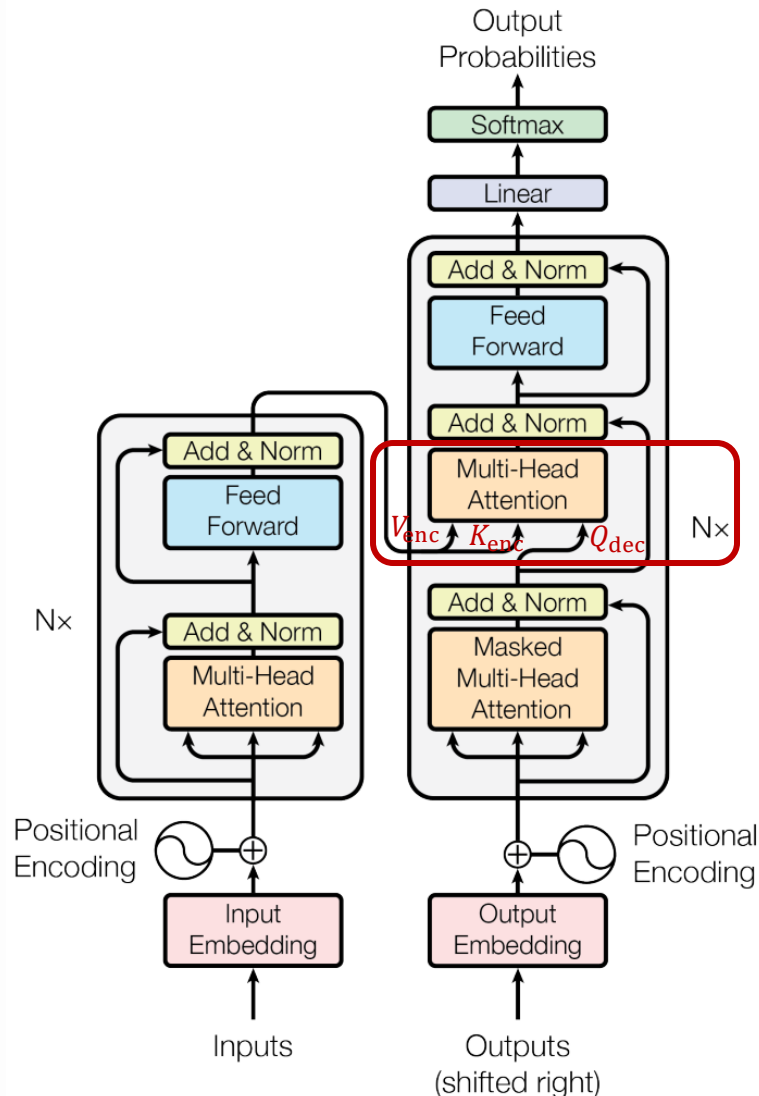
$$\text{softmax}(S + \mathbf{M})V = \begin{bmatrix} v_1 \\ c_{21}v_1 + c_{22}v_2 \\ c_{31}v_1 + c_{32}v_2 + c_{33}v_3 \\ c_{41}v_1 + c_{42}v_2 + c_{43}v_3 + c_{44}v_4 \end{bmatrix} = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix} = H_{\text{dec}}$$

Query (from decoder):  $Q_{\text{dec}} = H_{\text{dec}}W_Q$


Key (from encoder):  $K_{\text{enc}} = H_{\text{enc}}W_K$

Value (from encoder):  $V_{\text{enc}} = H_{\text{enc}}W_V$

$$\text{CrossAttention}(Q_{\text{dec}}, K_{\text{enc}}, V_{\text{enc}}) = \text{softmax}\left(\frac{Q_{\text{dec}}K_{\text{enc}}^T}{\sqrt{d_k}}\right)V_{\text{enc}}$$



# Decoder-only Architecture of Transformer



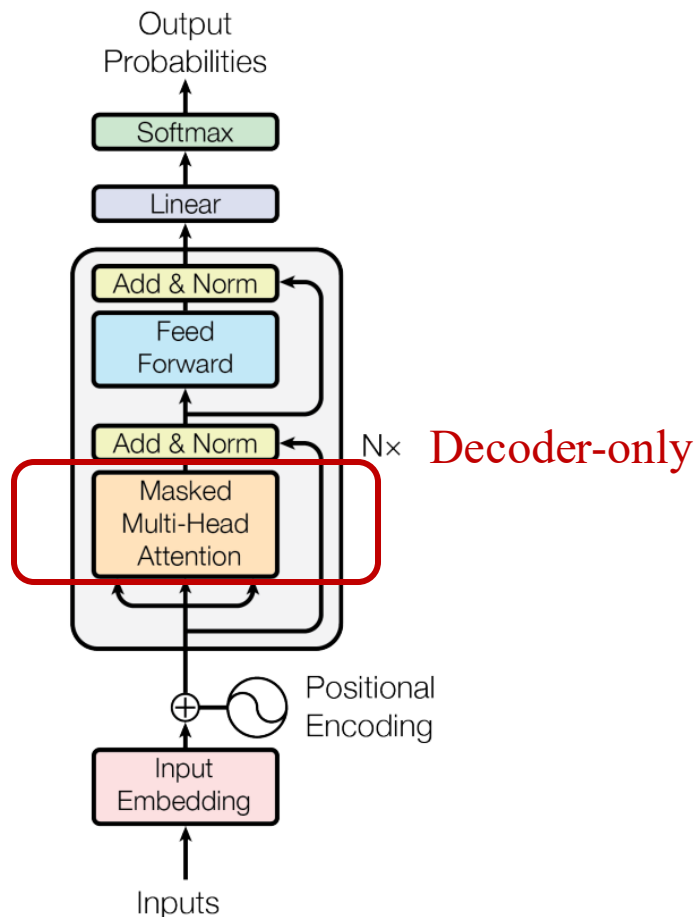
Encoder-decoder architecture

Example: vanilla Transformer

Decoder-only architecture

Example: GPT, LLaMA

# Decoder-only Architecture of Transformer




$$\text{MaskedAttention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}} + \mathbf{M}\right)V$$

$$\mathbf{M} = \begin{bmatrix} 0 & -\infty & -\infty & -\infty \\ 0 & 0 & -\infty & -\infty \\ 0 & 0 & 0 & -\infty \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

Advantages over encoder-decoder architecture

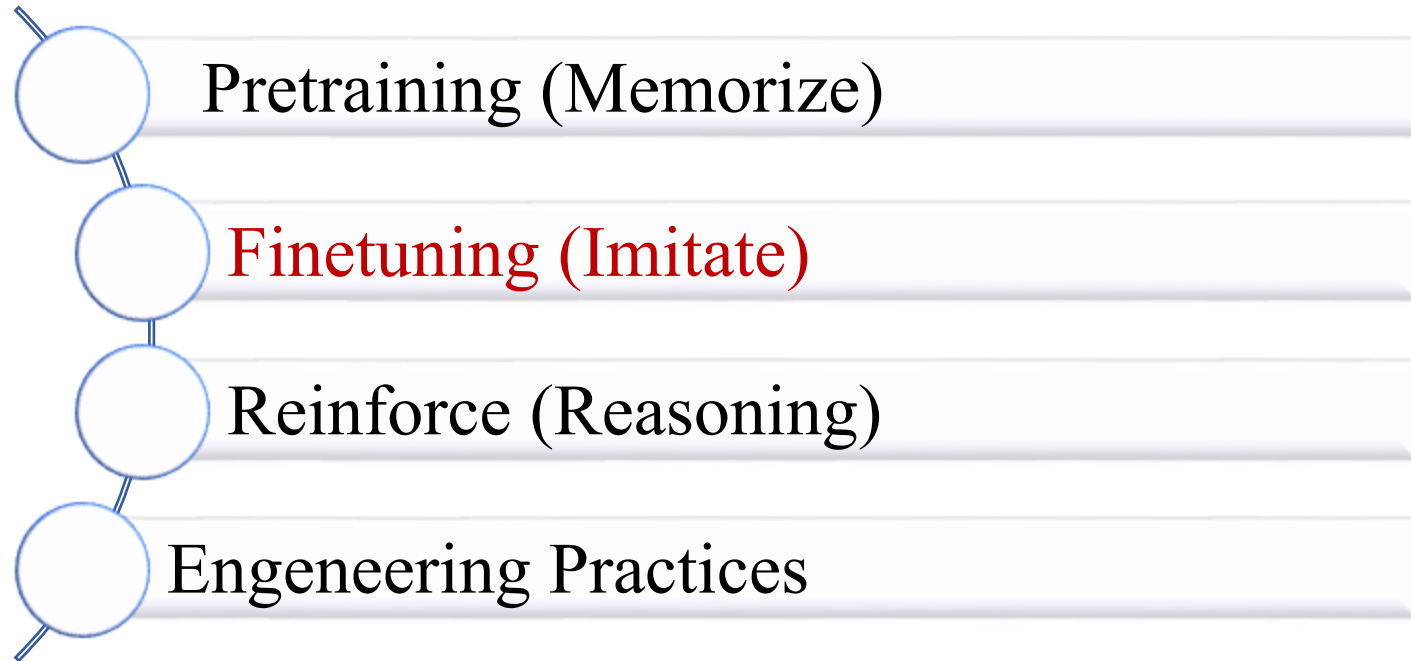
- Simplified architecture
- Appropriate for the generative tasks

# Performance of Decoder-only Transformer




■ Performance increases with the model size, but the increment is **limited**

# Outline



# Supervised FineTuning (SFT)



Prompt and response of Supervised FineTuning (SFT)

Prompt:  $[x_1, x_2, \dots, x_T]$

Response:  $[y_1, y_2, \dots, y_K]$ , **Provided by annotator**

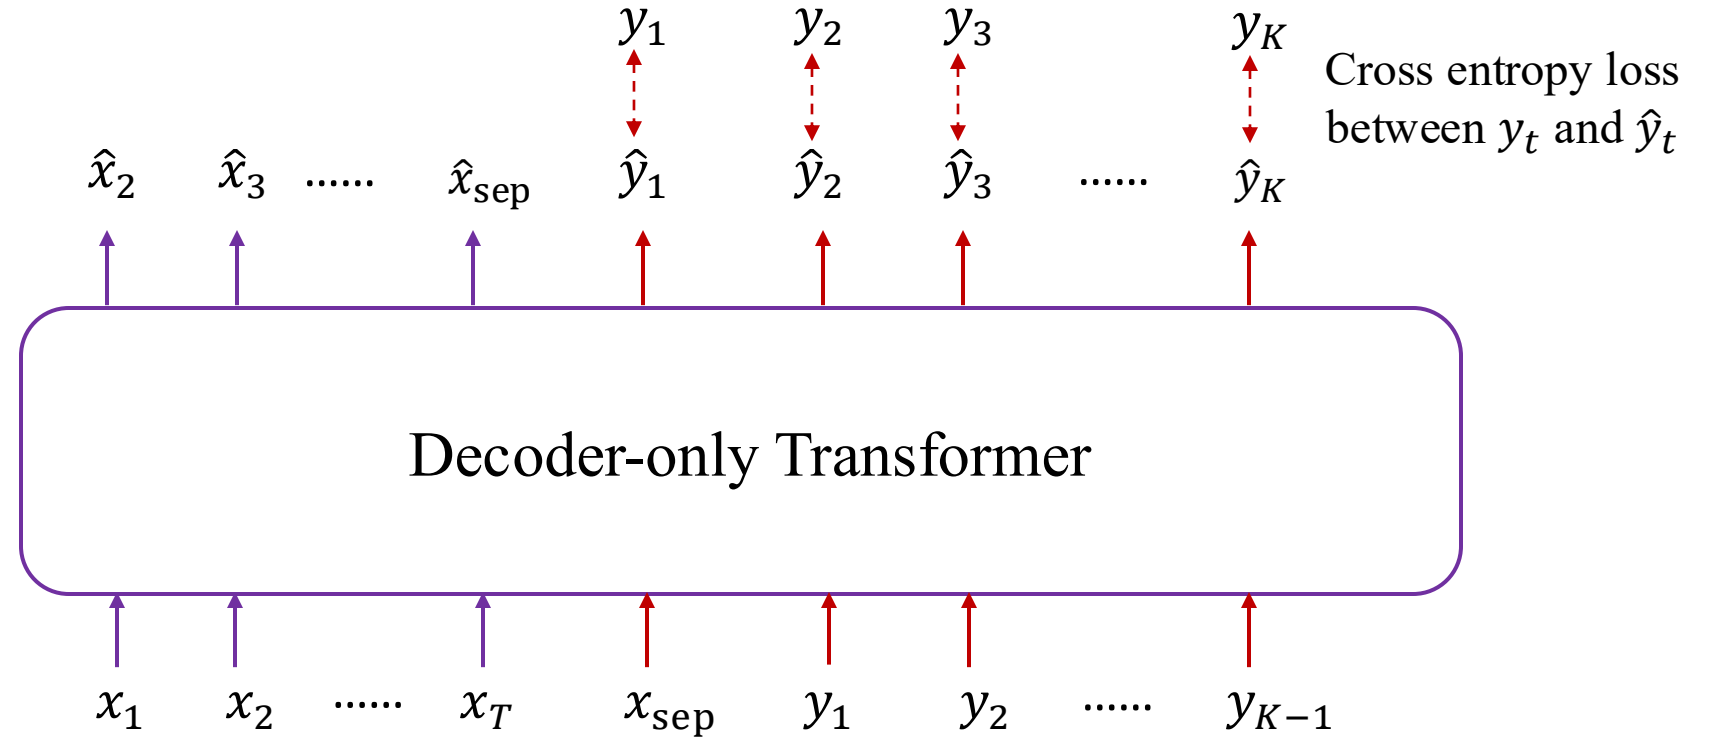
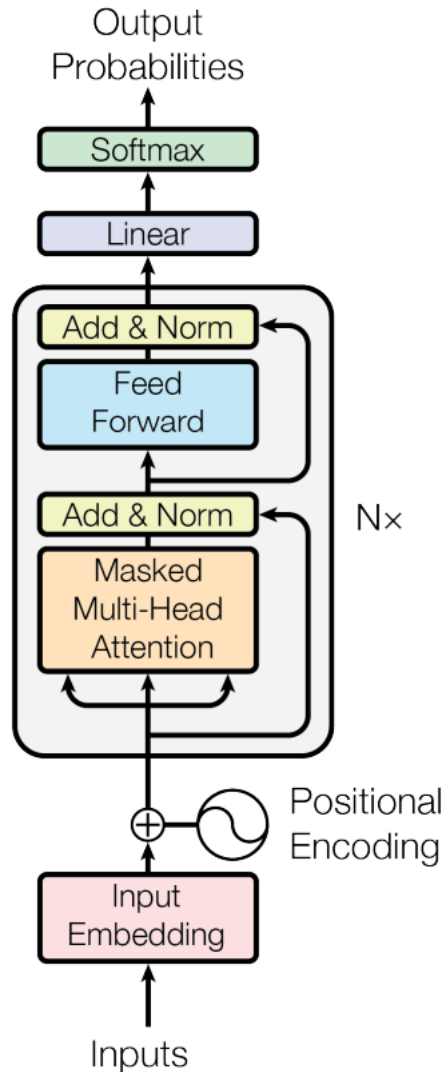
Training data of STF

$[x_1, x_2, \dots, x_T, < \text{sep} >, y_1, y_2, \dots, y_K]$

- Concatenate prompt and response to form training sample
- SFT is also an **autoregressive generative task** with the input consists of concatenated prompt and response




# Supervised FineTuning (SFT)



- Generate prompt and response as an autoregressive task, but calculate cross entropy loss only on the generated response
- Teacher forcing: feed input into the model despite the **correctness** of output

# Supervised FineTuning (SFT)



✓ Advantage: SFT can significantly increasing the performance

✗ Disadvantage: requires annotations (responses) from human

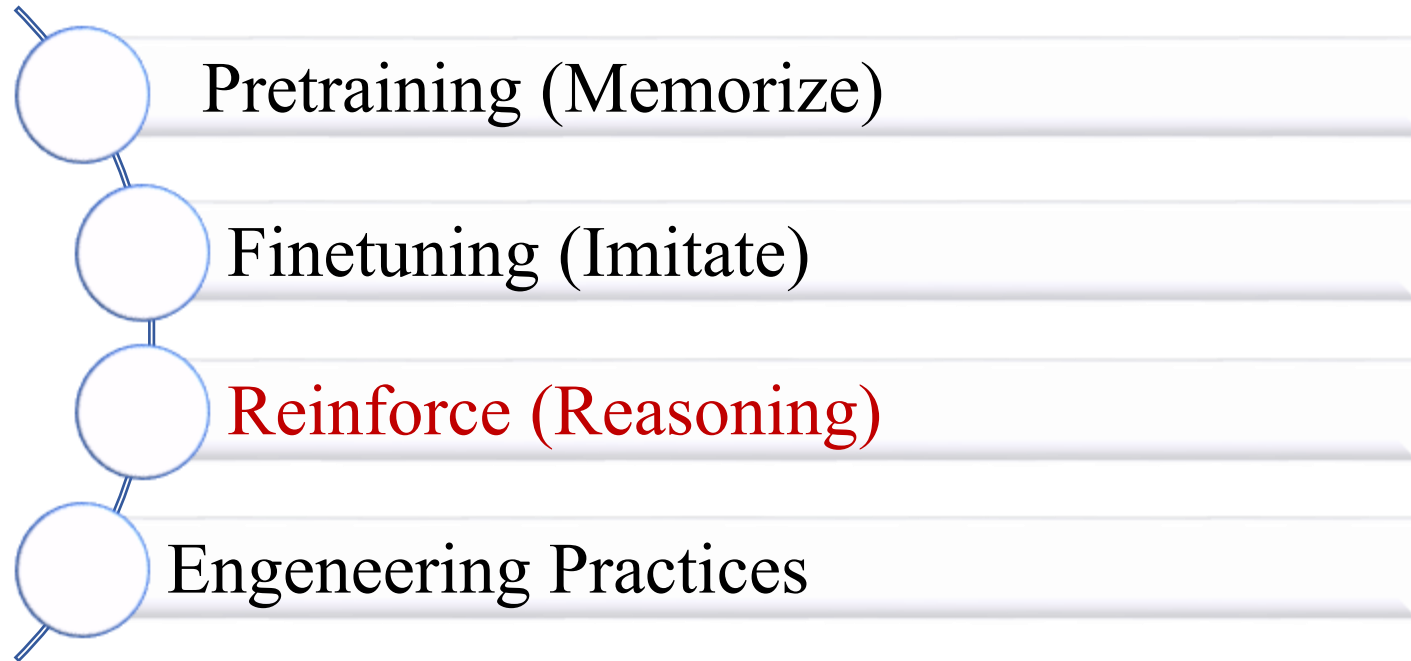
■ **Unsupervised** (self-supervised) pretraining: large scale unannotated dataset, but limited performance

■ **Supervised** finetuning: significant improvement in performance, but dataset is quite expensive

How to further improve the performance of LLM with manageable cost?

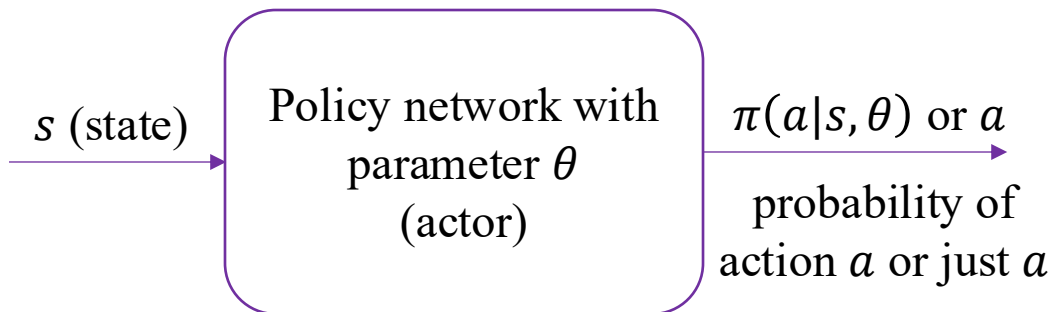
**Reinforcement Learning**

# Outline

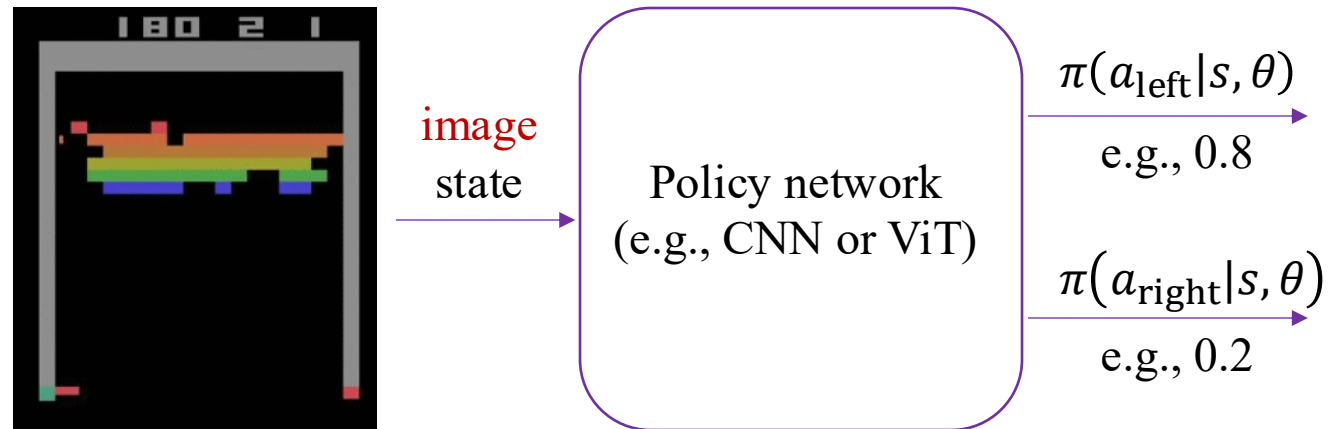


# Policy Gradient in RL


Policy gradient method




Example of Atari game



Example of LLM



# Reward Model in RLHF



- $x$  and  $\hat{y}$  are abbreviated as prompt and generated response, respectively


Two questions

- How to **measure** the quality of  $\hat{y}$  —————> **Reward model**
- How to **improve** the quality of  $\hat{y}$  —————> **Policy gradient**

# Reward Model in RLHF

■ How to **measure** the quality of  $\hat{y}$

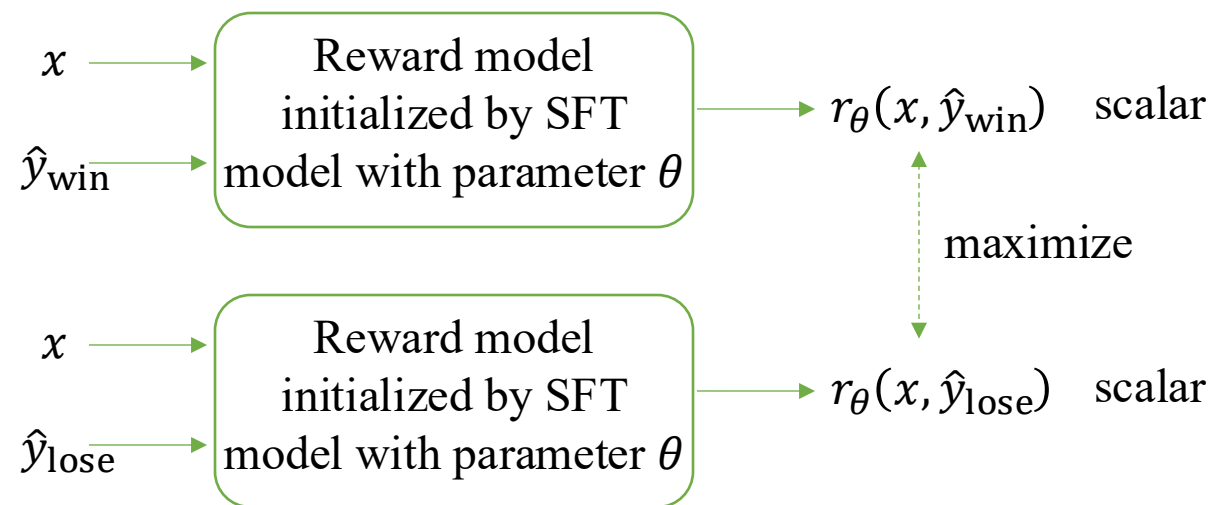
■ How to **improve** the quality of  $\hat{y}$



■ Generate  $M$  (e.g., 9) possible responses given the same prompt  $x$

■ Rank the quality of  $M$  responses by annotators

comparison between each two responses ( $\hat{y}_{\text{win}}$  and  $\hat{y}_{\text{lose}}$ ) is used as the signal to train a reward model



$$\text{maximize } r_{\theta}(x, \hat{y}_{\text{win}}) - r_{\theta}(x, \hat{y}_{\text{lose}})$$


$\sigma$ : Sigmoid function

$$\text{maximize } \log \left( \sigma \left( r_{\theta}(x, \hat{y}_{\text{win}}) - r_{\theta}(x, \hat{y}_{\text{lose}}) \right) \right)$$

$$\text{minimize } -\log \left( \sigma \left( r_{\theta}(x, \hat{y}_{\text{win}}) - r_{\theta}(x, \hat{y}_{\text{lose}}) \right) \right)$$

This is binary classification task

# Reward Model in RLHF




- Larger  $r_\theta(x, \hat{y})$  indicates  $\hat{y}$  is preferred
- Otherwise,  $\hat{y}$  is unpreferred

How to further **optimize the SFT** model to obtain high  $r_\theta(x, \hat{y})$  with powerful generative capability?

**Policy gradient method**

# Policy Gradient in RLHF



## Object 1

Maximize  $r_\theta(x, \hat{y})$ , where  $x$  and  $\hat{y}$  are drawn from the updated policy network ( $\pi^{RL}$ )

$$\text{Maximize } \mathbb{E}_{(x, \hat{y}) \sim D_{\pi^{RL}}} [r_\theta(x, \hat{y})]$$

## Object 2

Control the magnitude of policy network updates, avoid unstable training

$$\text{Maximize } \mathbb{E}_{(x, \hat{y}) \sim D_{\pi^{RL}}} \left[ -\beta \log \left( \frac{\pi^{RL}(\hat{y}|x)}{\pi^{SFT}(\hat{y}|x)} \right) \right]$$

$\pi^{RL}$ : updated policy network

$\pi^{SFT}$ : frozen policy network (or reference model) initialized by SFT


## Object 3 (optional)

Updated policy network ( $\pi^{RL}$ ) can still accomplish auto-regressive generative task, resulting **PPO-ptx**

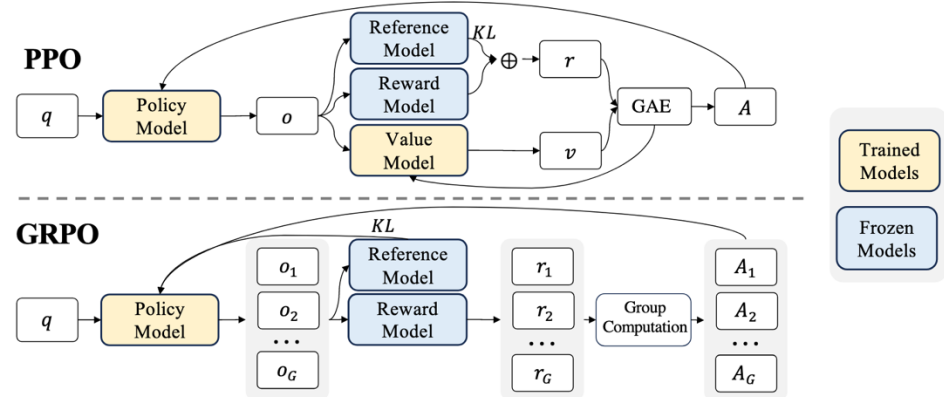
$$\text{Maximize } \mathbb{E}_{x \sim D_{\text{pretrain}}} [\gamma \log(\pi^{RL}(x))]$$




# RLHF



- Policy gradient significantly increasing the performance
- Potential of RL in LLM has not been fully explored

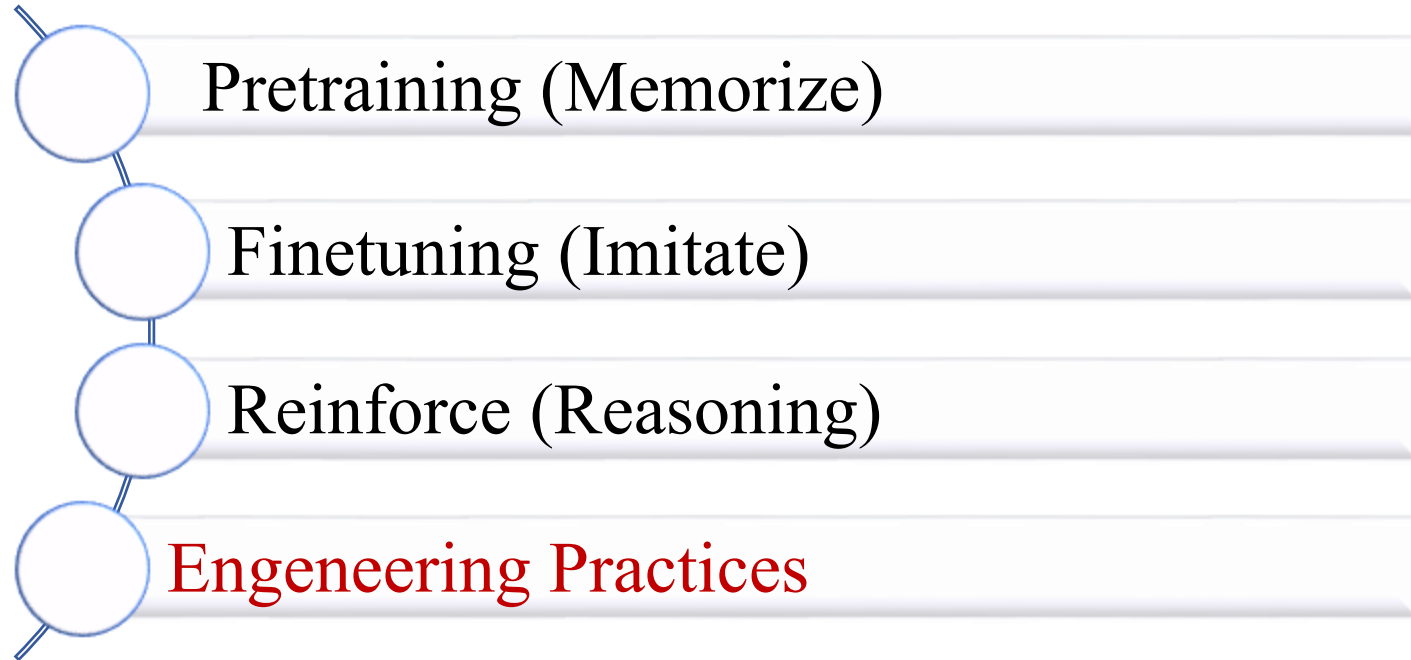


- Variant of PPO (GRPO) proposed by DeepSeek



- Vanilla PPO is enough (Open Reasoner Zero)

# Outline



# Important Library




## Transformers

Transformers provides thousands of pretrained models to perform tasks on different modalities such as text, vision, and audio.



TRL is a full stack library that provides a set of tools to train transformer language models with Reinforcement Learning, from the Supervised Fine-tuning (SFT), Reward Modeling (RM) to the Proximal Policy Optimization (PPO)



## Unsloth

Unsloth provides a collection of commonly used LLMs and accelerates the finetuning process by 2x, reduces memory usage by 70%, all while maintaining the same level of accuracy



## Datasets

Datasets provides dataloaders for numerous public datasets and efficient preprocessing methods

# Prepare Finetuning Dataset

“**Question**”: “一个1岁的孩子在夏季头皮出现多处小结节，长期不愈合， xxx”

“**CoT**”: “用中医的角度来看，出现小结节、再加上长期不愈合， xxx”

Chain of Thought (CoT), Optional

“**Response**”: “这是一种因湿热导致的疾病， xxx”

```
from datasets import load_dataset  
  
dataset = load_dataset("json", data_files="./data.json")
```

# Prepare Model

## ■ load model and tokenizer

```
from unsloth import FastLanguageModel
```

```
model, tokenizer = FastLanguageModel.from_pretrained(  
    model_name = "./DeepSeek-R1-Distill-Qwen-32B",  
    local_files_only=True,  
    max_seq_length = 4096,  
    load_in_4bit = True)
```

## ■ configure pretraining model

```
FastLanguageModel.for_training(model)
```

```
model = FastLanguageModel.get_peft_model(  
    model,  
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj"],  
    r=32, # LoRA related parameter  
    lora_alpha=16, # LoRA related parameter  
    use_gradient_checkpointing="unsloth")
```

# Prepare Trainer

## ■ load trainer

```
from trl import SFTTrainer
from transformers import TrainingArguments

trainer = SFTTrainer(
    model = model,
    tokenizer = tokenizer,
    train_dataset = dataset,
    dataset_text_field = "text", # 数据集字段名称
    max_seq_length = 4096,
    dataset_num_proc = 2, # 处理数据集进程数目
    args = TrainingArguments(
        per_device_train_batch_size = 2, # 每个GPU训练batch
        learning_rate = 2e-4, # 学习率
        optim = "adamw_8bit", # 使用8位AdamW优化器节省显存
        weight_decay = 0.01, # 正则化强度
        output_dir = "outputs", # 模型输出目录
        run_name = "medical-o1-sft-experiment", # 实验名称
    ),
)
```

`trainer.train()`

Take a break and let the LLM do the magic



**Thanks!**