



四川大學
SICHUAN UNIVERSITY

机器学习-第十五章 强化学习简介

教师：胡俊杰 副教授

邮箱：hujunjie@scu.edu.cn

1.凸优化概述

■ 优化问题的一般形式

$$\text{minimize } f_0(x)$$

待优化的目标函数

$$\begin{array}{l} \text{subject to } f_i(x) \leq 0, i = 1, 2, \dots, m \\ \text{(s.t.)} \end{array}$$

m 个不等式约束

$$h_i(x) = 0, i = 1, 2, \dots, p$$

p 个等式约束

1.凸优化概述

$$\text{minimize } f_0(x)$$

$$\text{s.t. } f_i(x) \leq 0, i = 1, 2, \dots, m$$

$$h_i(x) = 0, i = 1, 2, \dots, p$$

■ 拉格朗日函数

- 为每个约束指定一个拉格朗日乘子，以乘子为加权系数将约束增加到目标函数中

$$L(x, \lambda, v) = f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^p v_i h_i(x)$$

$$L(x, \lambda, v): R^n \times R_+^m \times R^p \rightarrow R$$

$$x \in R^n \quad \lambda \in R_+^m, v \in R^p$$

1.凸优化概述

■ 拉格朗日函数

$$L(x, \lambda, v) = f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^p v_i h_i(x)$$

$$\text{minimize } f_0(x)$$

$$\text{s.t. } f_i(x) \leq 0, i = 1, 2, \dots, m$$

$$h_i(x) = 0, i = 1, 2, \dots, p$$

■ 拉格朗日对偶函数

- 对拉格朗日函数 $L(x, \lambda, v)$ 中的 x 取下确界可定义拉格朗日对偶函数

$$g(\lambda, v) = \inf_{x \in R^n} L(x, \lambda, v) = \inf_{x \in R^n} \left(f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^p v_i h_i(x) \right)$$

- 拉格朗日对偶函数 $g(\lambda, v): R_+^m \times R^p \rightarrow R$

inf (infimum): 下确界, 数学分析中的概念, 小于等于集合中的所有成员的最大实数

$$\inf\{x \in R: 0 < x < 1\} = 0$$

sup (supremum): 上确界, 大于等于集合中所有成员的最小实数

$$\sup\{x \in R: 0 < x < 1\} = 1$$

1.凸优化概述

■ 拉格朗日对偶函数

$$g(\lambda, v) = \inf_{x \in \mathbb{R}^n} L(x, \lambda, v) = \inf_{x \in \mathbb{R}^n} \left(f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^p v_i h_i(x) \right)$$

- 拉格朗日对偶函数是凹函数，无论原问题是否为凸问题
- 拉格朗日对偶函数给出了原问题最优值的下界： $g(\lambda, v) \leq p^*$ ， p^* ：原问题 (primal problem) 的最优值 (optimal value)

1. 凸优化概述

■ 拉格朗日对偶函数

$$g(\lambda, v) = \inf_{x \in \mathbb{R}^n} L(x, \lambda, v) = \inf_{x \in \mathbb{R}^n} \left(f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{i=1}^p v_i h_i(x) \right)$$



从拉格朗日对偶函数获得的下界中，哪个是最优的？

当 $g(\lambda, v) = -\infty$ ，则其提供的下界无实际意义

■ 拉格朗日对偶问题

$$\max_{\lambda \geq 0, v} g(\lambda, v) = \max_{\lambda \geq 0, v} \inf_{x \in \mathbb{R}^n} L(x, \lambda, v)$$

- 假设拉格朗日对偶问题 (dual problem) 的最优值为 d^*

2.支持向量机概述

- 给定训练数据集 $D = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$, 其中


$$x_i \in R^n, y_i \in \{+1, -1\}, i = 1, 2, \dots, N$$

假设超平面 (w, b) 将线性可分数据集正确分类, 即对于 $(x_i, y_i) \in D$, 若 $y_i = +1$, 则 $w^T x_i + b > 0$; 若 $y_i = -1$, 则 $w^T x_i + b < 0$, 令

$$\begin{cases} w^T x_i + b \geq +1, y_i = +1 \\ w^T x_i + b \leq -1, y_i = -1 \end{cases}$$

- 将以上两个方程合并, 可得: $y_i(w^T x_i + b) \geq 1$


- $y_i(w^T x_i + b) \geq 1$ 为**约束条件**, 即要求决策超平面 (w, b) 将所有样本分类正确



2.支持向量机概述

- 支持向量到超平面的距离可以写为: $d = \frac{|w^T x + b|}{\|w\|} = \frac{1}{\|w\|}$
- 两类（正类和负类）支持向量到超平面的距离之和为 $\gamma = \frac{2}{\|w\|}$ ，也被称为“间隔”，即**优化目标**

$$\begin{aligned} & \max_{w,b} \frac{2}{\|w\|} \\ & s.t. y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N \end{aligned}$$



2.支持向量机概述

$$\max_{w,b} \frac{2}{\|w\|}$$

$$s.t. y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$



$$\min_{w,b} \frac{1}{2} \|w\|^2 \quad \text{方便后续求导}$$

$$s.t. y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$

■ 即支持向量机 (SVM) 的基本形式

3.支持向量机求解

原优化问题:

$$\min_{w,b} \frac{1}{2} \|w\|^2$$
$$s.t. y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$

拉格朗日函数: $L(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^N \alpha_i y_i (w^T x_i + b) + \sum_{i=1}^N \alpha_i$

- 线性可分SVM满足强对偶条件, 即 $d^* = p^*$, 原问题最优值等于对偶问题最优值
- 通过解对偶问题, 得到最优解 $\alpha^* = (\alpha_1^*, \alpha_2^*, \dots, \alpha_N^*)$, 得到原问题的最优解 (w^*, b^*)

3.支持向量机求解

线性可分支持向量机学习算法

第1步：根据原始优化问题，写出拉格朗日函数

$$L(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^N \alpha_i y_i (w^T x_i + b) + \sum_{i=1}^N \alpha_i$$

第2步：求 $\min_{w,b} L(w, b, \alpha)$ ，并代入 $L(w, b, \alpha)$

$$\min_{w,b} L(w, b, \alpha) = -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (x_i \cdot x_j) + \sum_{i=1}^N \alpha_i$$

第3步：求解拉格朗日对偶问题，即

$$\max_{\alpha} \min_{w,b} L(w, b, \alpha) = \max_{\alpha} -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (x_i \cdot x_j) + \sum_{i=1}^N \alpha_i \quad s.t. \sum_{i=1}^N \alpha_i y_i = 0$$
$$\alpha_i \geq 0, i = 1, 2, \dots, N$$

求得最优解 $\alpha^* = (\alpha_1^*, \alpha_2^*, \dots, \alpha_N^*)$

第4步：根据KKT条件可得原优化问题最优解 w^* 和 b^*


$$w^* = \sum_{i=1}^N \alpha_i y_i x_i \quad b^* = y_j - \sum_{i=1}^N \alpha_i^* y_i (x_i \cdot x_j)$$

第5步：构建决策超平面以及分类决策函数

决策超平面： $w^{*T} x + b^* = 0$

决策函数： $f(x) = \text{sign}(w^{*T} x + b^*)$


1.线性支持向量机




$$\min_{w,b} \frac{1}{2} \|w\|^2$$

$$s. t. y_i (w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$

- 以上假设训练样本是线性可分的，即存在一个超平面将两类样本完全分开
- 然而现实情况很复杂（如标签错误），应允许支持向量机对部分样本出错



1.线性支持向量机



■ 允许一些样本不满足约束 $y_i(w^T x_i + b) \geq 1$

■ 在最大化间隔的同时，不满足约束的样本应尽可能少，目标如下


$$\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N l_{0/1}(y_i(w^T x_i + b) - 1)$$

● $C > 0$ 为人为指定的惩罚参数， $l_{0/1}$ 代表0/1损失函数

$$l_{0/1} = \begin{cases} 1, & \text{if } z < 0 \\ 0, & \text{otherwise} \end{cases}$$

● 常量 C 允许部分样本不满足约束 $y_i(w^T x_i + b) \geq 1$


1.线性支持向量机



$$\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N l_{0/1}(y_i(w^T x_i + b) - 1)$$

- $l_{0/1}$ 非凸、不连续，以上目标函数不易求解。用数学性质较好的 “代理损失函数” (surrogate loss function)，替代 $l_{0/1}$
- 常用代理损失函数：
 - hinge损失函数 (hinge loss) : $l_{hinge}(z) = \max(0, 1 - z)$
 - 指数损失函数 (exponential loss) : $l_{exp}(z) = e^{-z}$
 - 对率损失函数 (logistic loss) : $l_{log}(z) = \log(1 + e^{-z})$

1.线性支持向量机



$$\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N l_{0/1}(y_i(w^T x_i + b) - 1)$$

$l_{hinge}(z) = \max(0, 1 - z)$ 代替 $l_{0/1}$

$$\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \max(0, 1 - y_i(w^T x_i + b))$$

$\xi_i = \max(0, 1 - y_i(w^T x_i + b))$
(/ksaɪ/)

$$\begin{aligned} \min_{w,b,\xi_i} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & y_i(w^T x_i + b) \geq 1 - \xi_i \\ & \xi_i \geq 0, i = 1, 2, \dots, N \end{aligned}$$

优化目标及约束

1.线性支持向量机

线性支持向量机的拉格朗日对偶问题

$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (x_i \cdot x_j) - \sum_{i=1}^N \alpha_i \\ \text{s.t.} \quad & \sum_{i=1}^N \alpha_i y_i = 0 \\ & 0 \leq \alpha_i \leq C, i = 1, 2, \dots, N \end{aligned}$$

C 为人为指定的
惩罚参数

线性可分支持向量机的拉格朗日对偶问题


$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (x_i \cdot x_j) - \sum_{i=1}^N \alpha_i \\ \text{s.t.} \quad & \sum_{i=1}^N \alpha_i y_i = 0 \\ & 0 \leq \alpha_i, i = 1, 2, \dots, N \end{aligned}$$

■ 设 $\alpha^* = (\alpha_1^*, \alpha_2^*, \dots, \alpha_N^*)$ 为拉格朗日对偶问题最优解，则原问题最优解 w^* 和 b^* 如下


$$w^* = \sum_{i=1}^N \alpha_i^* y_i x_i \quad b^* = y_j - \sum_{i=1}^N \alpha_i^* y_i (x_i \cdot x_j)$$

由KKT中的互补松弛条件可得

3.非线性支持向量机



线性不可分



高维下可分

- 无法找到一个超平面将两类样本分开
- 解决思路：对输入 x 作用非线性变换 ϕ ，将 x 从原始空间映射至高维空间，使得 $\phi(x)$ 可分

3.非线性支持向量机

线性可分支持向量机

$$\min_{w,b} \frac{1}{2} \|w\|^2$$

$$s.t. y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$

拉格朗日对偶问题如下:

$$\max_{\alpha} -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (x_i \cdot x_j) + \sum_{i=1}^N \alpha_i$$

$$s.t. \sum_{i=1}^N \alpha_i y_i = 0$$

$$\alpha_i \geq 0, i = 1, 2, \dots, N$$

非线性支持向量机, 优化目标及约束如下:

$$\min_{w,b} \frac{1}{2} \|w\|^2$$

$$s.t. y_i(w^T \phi(x_i) + b) \geq 1, i = 1, 2, \dots, N$$

拉格朗日对偶问题如下:

$$\max_{\alpha} -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (\phi(x_i) \cdot \phi(x_j)) + \sum_{i=1}^N \alpha_i$$

$$s.t. \sum_{i=1}^N \alpha_i y_i = 0$$

$$0 \leq \alpha_i \leq C, i = 1, 2, \dots, N$$

3.非线性支持向量机

非线性支持向量机，优化目标及约束如下：

$$\min_{w,b} \frac{1}{2} \|w\|^2$$

$$s. t. y_i (w^T \phi(x_i) + b) \geq 1, i = 1, 2, \dots, N$$

拉格朗日对偶问题如下：

$$\max_{\alpha} -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j (\phi(x_i) \cdot \phi(x_j)) + \sum_{i=1}^N \alpha_i$$

$$s. t. \sum_{i=1}^N \alpha_i y_i = 0$$

$$0 \leq \alpha_i \leq C, i = 1, 2, \dots, N$$

- 直接定义非线性映射函数 ϕ 较困难
- 引入核函数 $K(x_i, x_j) = \phi(x_i) \cdot \phi(x_j)$ ， x_i 与 x_j 在特征空间中的内积等于在原始输入空间通过函数 $K(x_i, x_j)$ 计算的结果，从而避开定义 ϕ

3.非线性支持向量机

线性可分支持向量机

$$w^* = \sum_{i=1}^N \alpha_i^* y_i x_i \quad b^* = y_j - \sum_{i=1}^N \alpha_i^* y_i (x_i \cdot x_j)$$

超平面: $\sum_{i=1}^N \alpha_i^* y_i (x \cdot x_i) + b^* = 0$

决策函数: $f(x) = \text{sign} \left(\sum_{i=1}^N \alpha_i^* y_i (x \cdot x_i) + b^* \right)$

非线性支持向量机

$$w^* = \sum_{i=1}^N \alpha_i^* y_i \phi(x_i) \quad b^* = y_j - \sum_{i=1}^N \alpha_i^* y_i \phi(x_i) \cdot \phi(x_j)$$

$K(x_i, x_j)$

超平面: $\sum_{i=1}^N \alpha_i^* y_i \phi(x) \cdot \phi(x_i) + b^* = 0$

$K(x, x_i)$

决策函数: $f(x) = \text{sign} \left(\sum_{i=1}^N \alpha_i^* y_i \phi(x) \cdot \phi(x_i) + b^* \right)$

$K(x, x_i)$

3.序列最小优化算法

拉格朗日对偶问题：

$$\max_{\alpha} -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j K(x_i, x_j) + \sum_{i=1}^N \alpha_i$$

$$s.t. \sum_{i=1}^N \alpha_i y_i = 0$$

$$0 \leq \alpha_i \leq C, i = 1, 2, \dots, N$$

■ 如何求解 α^* ？

■ 序列最小优化算法（Sequential Minimal Optimization）：每次选择 α_i 和 α_j ，并固定其他参数，求 α_i 和 α_j 的极值，直至收敛。求解思路如下：

- 选择一对需要更新的变量 α_i 和 α_j
- 固定 α_i 和 α_j 之外的参数，求 ∇_{α_i} 和 ∇_{α_j}

■ SMO算法的两个主要部分：

- 求解两个变量 α_i 和 α_j 的解析方法
- 选择待优化变量 α_i 和 α_j 的方法

本章目录

01 面向决策任务的人工智能

02 强化学习的基础概念

03 Bellman equation



两种人工智能类型

□ 预测型任务

- 根据数据预测所需输出 (有监督学习)
- 数据降维或得到表达 (无监督学习)



□ 决策型任务

- 在动态环境中采取行动 (强化学习)
 - 转变到新的状态
 - 获得即时奖励
 - 随着时间的推移最大化累计奖励



两种人工智能类型

- ❑ 决策下达到环境中，直接改变环境
 - 未来发展随之改变
- ❑ 预测仅产生信号，不考虑环境的改变
 - 不需要考虑预测的信号是否用、怎么用

智能医学

- 决策
 - 医生或者人工智能模型直接给病人下达治疗方案
- 预测
 - 人工智能模型告诉医生关于病人可能的得病预测，医生综合各方面判断给病人下达治疗方案



两种人工智能类型

□ 序贯决策

- 决策者序贯地做出一个个决策，并接续看到新的观测，直到最终任务结束。




绝大多数序贯决策问题，可以用**强化学习**来建模

两种人工智能类型

□ 序贯决策

- 决策者序贯地做出一个个决策，并接续看到新的观测，直到最终任务结束。



本章目录

- 01** 面向决策任务的人工智能
- 02** 强化学习的基础概念
- 03** Bellman equation



强化学习定义

有监督、无监督学习

Model



Fixed Data

强化学习


Agent

(Decision maker)



Dynamic Environment

Basic Concepts of RL



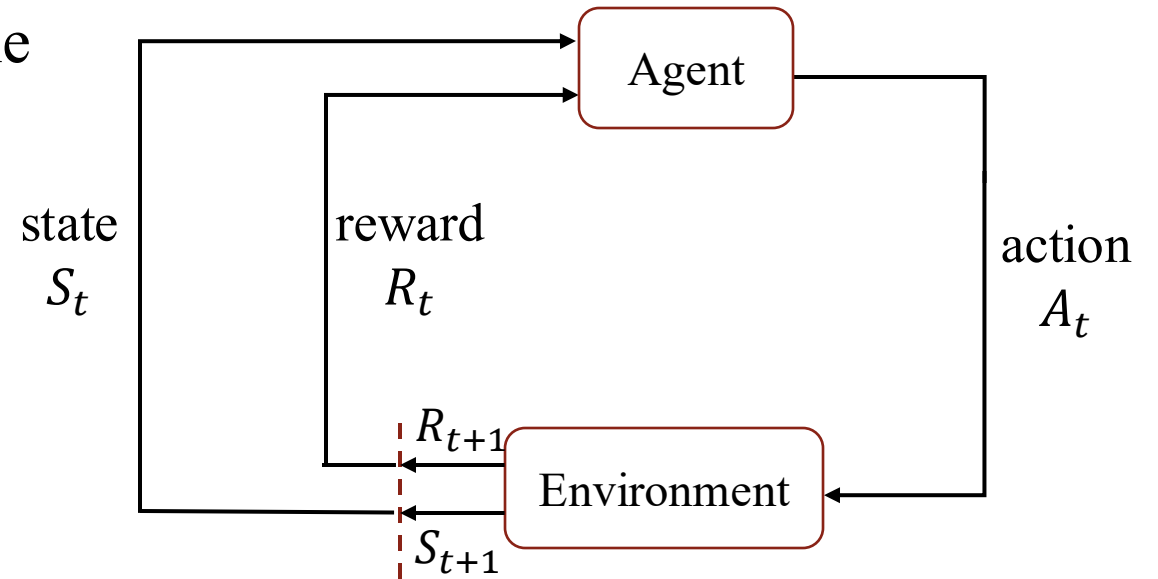
S_t, A_t, R_t are random variables.

Basic Concepts of RL : State

State: A state is a representation of the environment at a specific time.

State Space: $S = \{s_1, \dots, s_n\}$

State $S_t = s$

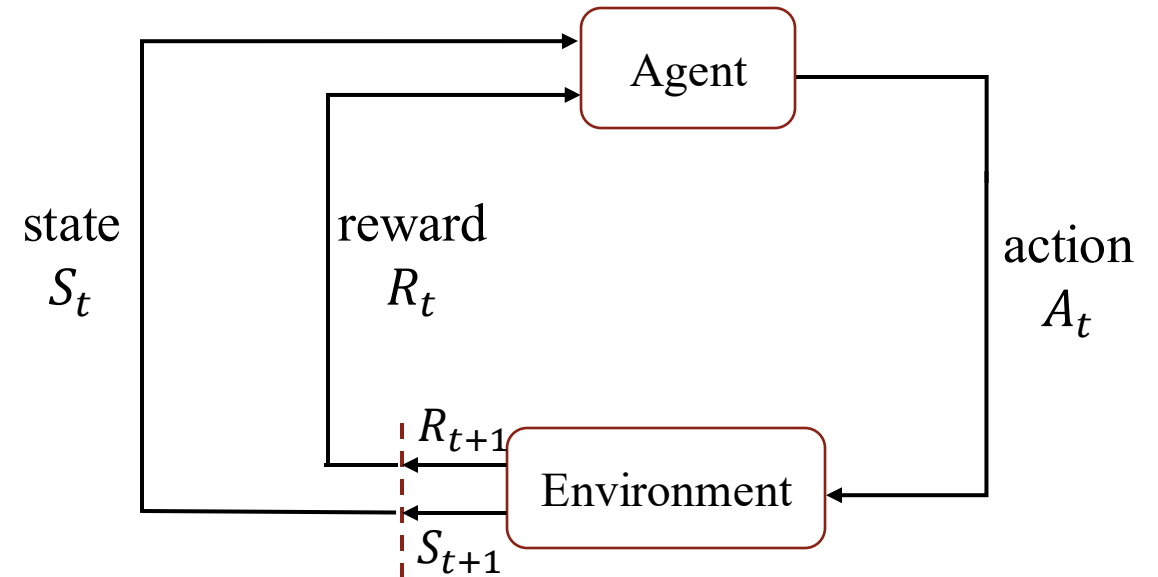


Basic Concepts of RL : Action

Action: An action is a decision or move made by the agent that affects the environment.

Action Space: $A = \{a_1, \dots, a_n\}$

Action $A_t = a$



Basic Concepts of RL : State Transition

State transition: A state transition describes how the environment moves from one state s to another state s' after the agent takes an action a .

$$s \xrightarrow[p(s'|s, a)]{a} s'$$

Basic Concepts of RL : Reward



Reward: a scalar signal that tells the agent how good or bad its action is in a given state.

The reward is typically defined as a function of the state or the state-action pair.


$r(s, a)$ with probability $p(r|s, a)$

Reward $R_t = r(s, a) = r$

Reward Space: $R = \{r_1, \dots, r_n\}$



Basic Concepts of RL : Return



Relationship between Return and Reward

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$


$$G_t = R_{t+1} + \gamma G_{t+1}$$

$\gamma \in [0,1]$: A hyperparameter that controls the trade-off between short-term and long-term rewards


Basic Concepts of RL : Policy

Policy: A policy tells the agent which actions to take at every state.


$$\pi(a|s) = p(a|s)$$




Action space: $A = \{a_1, \dots, a_n\}$



(a) States




(b) Actions



	a_1 (upward)	a_2 (rightward)	a_3 (downward)	a_4 (leftward)	a_5 (still)
s_1	0	0.5	0.5	0	0
s_2	0	0	1	0	0
s_3	0	0	0	1	0
s_4	0	1	0	0	0
s_5	0	0	1	0	0
s_6	0	0	1	0	0
s_7	0	1	0	0	0
s_8	0	1	0	0	0
s_9	0	0	0	0	1

Ultimate goal of RL: Find the **optimal policy**

Basic Concepts of RL : Markov Decision Processes (MDP)



Markov property: The next state s_{t+1} and reward r_{t+1} depend only on the current state s_t and action a_t , not on the full history.

$$S_t \xrightarrow{A_t} S_{t+1}, R_{t+1} \xrightarrow{A_{t+1}} S_{t+2}, R_{t+2} \xrightarrow{A_{t+2}} S_{t+3}, R_{t+3} \dots$$

$$p(s_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(s_{t+1}|s_t, a_t)$$

$$p(r_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = p(r_{t+1}|s_t, a_t)$$

Basic Concepts of RL : State Value

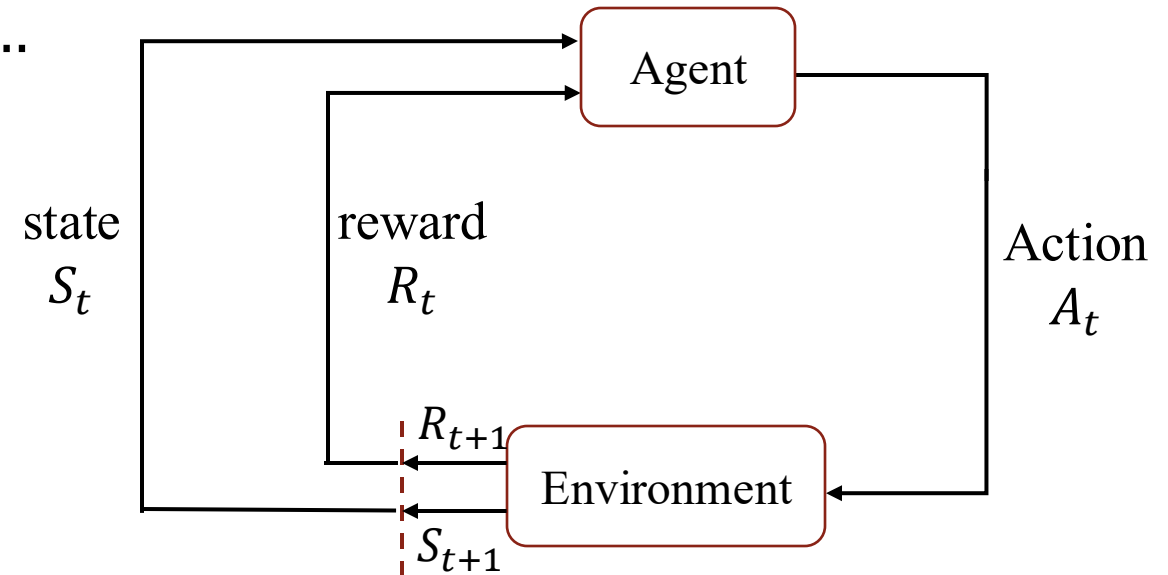
$$S_t \xrightarrow{A_t} S_{t+1}, R_{t+1} \xrightarrow{A_{t+1}} S_{t+2}, R_{t+2} \xrightarrow{A_{t+2}} S_{t+3}, R_{t+3} \dots \dots$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

State value $v_\pi(s)$: expected return (cumulative future reward) an agent can get by starting in state s and following policy π

State value measures how good it is for the agent in a state follows a certain policy

$$v_\pi(s) = E[G_t | S_t = s]$$




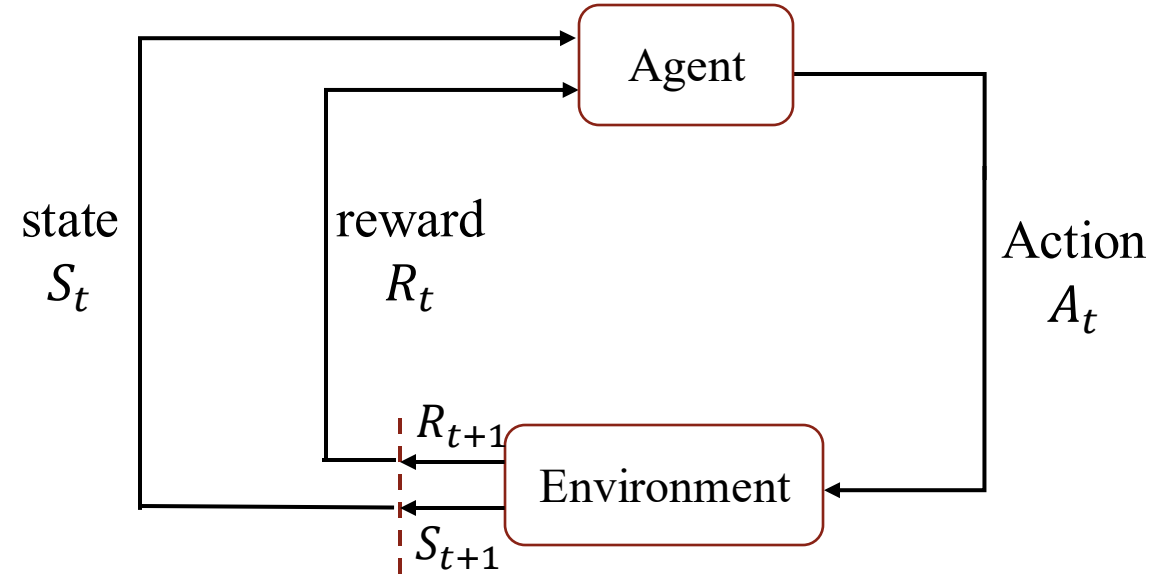
How to measure the effectiveness of action in state value?

Basic Concepts of RL : Action Value

Action value $q_\pi(s, a)$: expected return starting from state s , **taking action a** , and then following policy π

It tells how good it is to take action a in state s under policy π

$$q_\pi(s, a) = E[G_t | S_t = s, A_t = a]$$




Basic Concepts of RL : Relationship between State Value and Action Value

$$v_{\pi}(s) = E[G_t | S_t = s]$$

$$q_{\pi}(s, a) = E[G_t | S_t = s, A_t = a]$$

Law of total expectation
全期望公式

$$\mathbb{E}[X] = \sum_a \mathbb{E}[X | A = a] p(a)$$



$$v_{\pi}(s) = E[G_t | S_t = s] = \sum_{a \in A} E[G_t | S_t = s, A_t = a] \pi(a|s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a)$$

本章目录

01 面向决策任务的人工智能

02 强化学习的基础概念

03 Bellman equation



Bellman Equation

State value: $v_{\pi}(s) = E[G_t | S_t = s]$

Action value: $q_{\pi}(s, a) = E[G_t | S_t = s, A_t = a]$

Relationship between state value and action value:

$$v_{\pi}(s) = \sum_{a \in A} E[G_t | S_t = s, A_t = a] \pi(a|s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a)$$

How to solve the above equation to obtain state value and action value?

Bellman equation!

Bellman Equation

$$v_{\pi}(s) = E[G_t | S_t = s]$$

$$= E[R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= E[R_{t+1} | S_t = s] + \gamma E[G_{t+1} | S_t = s]$$

mean of immediate rewards

mean of future rewards

$$v_{\pi}(s) = E[G_t | S_t = s]$$

$$G_t = R_{t+1} + \gamma G_{t+1}$$

$$E[R_{t+1} | S_t = s] = \sum_{a \in A} E[R_{t+1} | S_t = s, A_t = a] p(a|s) = \sum_{a \in A} \pi(a|s) \sum_{r \in R} r(s, a) p(r|s, a) = \sum_{a \in A} \pi(a|s) \sum_{r \in R} r p(r|s, a)$$

$$E[G_{t+1} | S_t = s] = \sum_{s' \in S} E[G_{t+1} | S_t = s, S_{t+1} = s'] p(s'|s) = \sum_{s' \in S} E[G_{t+1} | S_{t+1} = s'] p(s'|s) = \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

Markov property

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

Bellman Equation

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

$$p(s'|s) = \sum_{a \in A} p(s'|s, a) \pi(a|s)$$

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \left[\sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_{\pi}(s') \right]$$

Bellman Equation


From action value to state value: $v_{\pi}(s) = \sum_{a \in A} E[G_t | S_t = s, A_t = a] \pi(a|s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a)$

From state value to action value: $v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \left(\sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s') \right)$

$$q_{\pi}(s, a) = \sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

Bellman Equation

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) \sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$


$$r_{\pi}(s)$$

$$v_{\pi}(s) = r_{\pi}(s) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

Vector form

$$v_{\pi} = r_{\pi} + \gamma P_{\pi} v_{\pi}$$

$$v_{\pi} = \begin{bmatrix} v_{\pi}(s_1) \\ v_{\pi}(s_2) \\ \vdots \\ v_{\pi}(s_n) \end{bmatrix} \quad r_{\pi} = \begin{bmatrix} r_{\pi}(s_1) \\ r_{\pi}(s_2) \\ \vdots \\ r_{\pi}(s_n) \end{bmatrix}$$

$$P_{\pi} = \begin{bmatrix} p(s_1|s_1) & \cdots & p(s_n|s_1) \\ \vdots & \ddots & \vdots \\ p(s_1|s_n) & \cdots & p(s_n|s_n) \end{bmatrix}$$

$$\sum_{j=1}^n p(s_j|s_i) = 1$$

Solution of Bellman Equation

Vector form

$$v_{\pi} = r_{\pi} + \gamma P_{\pi} v_{\pi}$$




$$v_{\pi} = [I - \gamma P_{\pi}]^{-1} r_{\pi}$$

$$v_{\pi} = \begin{bmatrix} v_{\pi}(s_1) \\ v_{\pi}(s_2) \\ \vdots \\ v_{\pi}(s_n) \end{bmatrix} \quad r_{\pi} = \begin{bmatrix} r_{\pi}(s_1) \\ r_{\pi}(s_2) \\ \vdots \\ r_{\pi}(s_n) \end{bmatrix}$$

$$P_{\pi} = \begin{bmatrix} p(s_1|s_1) & \cdots & p(s_n|s_1) \\ \vdots & \ddots & \vdots \\ p(s_1|s_n) & \cdots & p(s_n|s_n) \end{bmatrix}$$

$$\sum_{j=1}^n p(s_j|s_i) = 1$$

Example of Computing Action Value and State Value



$$q_{\pi}(s, a) = \sum_{r \in R} r p(r|s, a) + \gamma \sum_{s' \in S} p(s'|s) v_{\pi}(s')$$

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a)$$

Considering the a_2 and a_3 in s_1

$$q_{\pi}(s_1, a_2) = -1 + \gamma v_{\pi}(s_2)$$

$$q_{\pi}(s_1, a_3) = 0 + \gamma v_{\pi}(s_3)$$

$$\begin{aligned} v_{\pi}(s_1) &= 0.5 q_{\pi}(s_1, a_2) + 0.5 q_{\pi}(s_1, a_3) \\ &= 0.5[-1 + \gamma v_{\pi}(s_2)] + 0.5[-1 + \gamma v_{\pi}(s_3)] \end{aligned}$$

- ◇ If the agent attempts to exit the boundary, let $r_{\text{boundary}} = -1$.
- ◇ If the agent attempts to enter a forbidden cell, let $r_{\text{forbidden}} = -1$.
- ◇ If the agent reaches the target state, let $r_{\text{target}} = +1$.
- ◇ Otherwise, the agent obtains a reward of $r_{\text{other}} = 0$.

Value Iteration to find the optimal policy

model-based method

Initialization: The probability models $p(r|s, a)$ and $p(s'|s, a)$ for all (s, a) are known.
Initial guess v_0 .

Goal: Search for the optimal state value and an optimal policy for solving the Bellman optimality equation.

While v_k has not converged in the sense that $\|v_k - v_{k-1}\|$ is greater than a predefined small threshold, for the k th iteration, do

For every state $s \in \mathcal{S}$, do

For every action $a \in \mathcal{A}(s)$, do

q-value: $q_k(s, a) = \sum_r p(r|s, a)r + \gamma \sum_{s'} p(s'|s, a)v_k(s')$

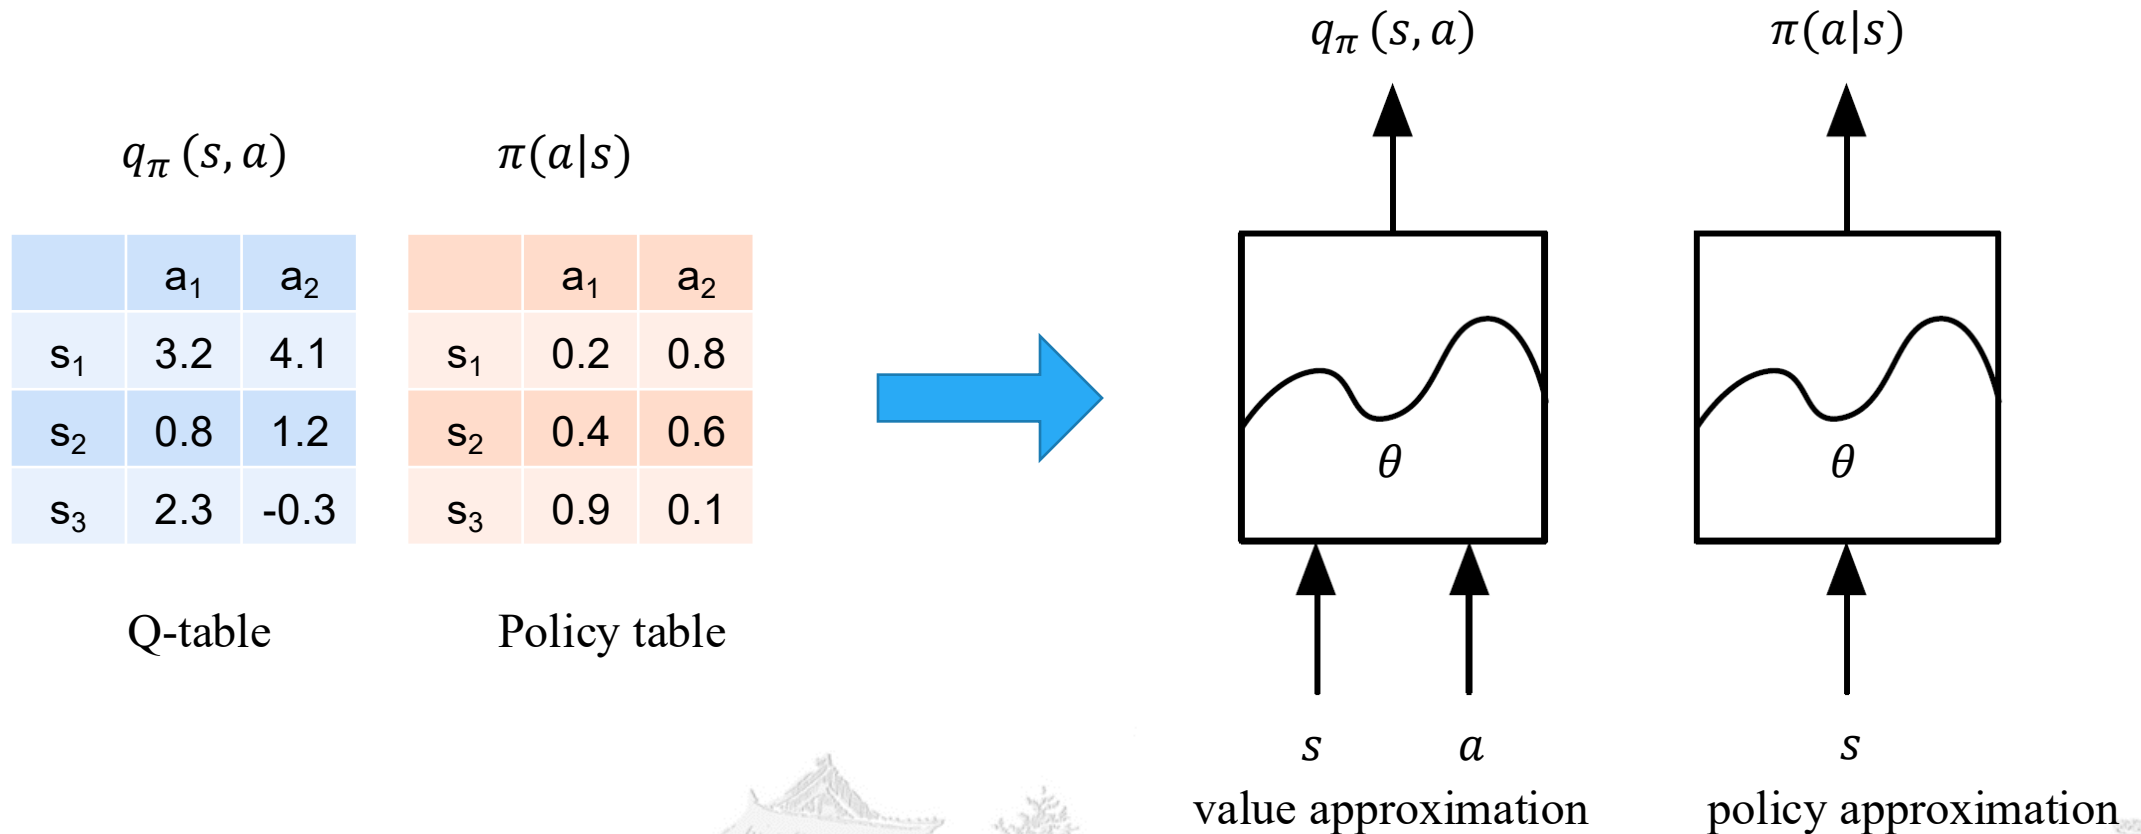
Maximum action value: $a_k^*(s) = \arg \max_a q_k(s, a)$

Policy update: $\pi_{k+1}(a|s) = 1$ if $a = a_k^*$, and $\pi_{k+1}(a|s) = 0$ otherwise

Value update: $v_{k+1}(s) = \max_a q_k(s, a)$



Q-table



Deep reinforcement learning: use deep neural networks to approximate the $q_{\pi}(s, a)$ and $\pi(a|s)$

Deep Reinforcement Learning

- 2012, AlexNet is proposed
- 2013, the first deep reinforcement learning paper was presented at the NIPS 2013

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih Koray Kavukcuoglu David Silver Alex Graves Ioannis Antonoglou

Daan Wierstra Martin Riedmiller

DeepMind Technologies

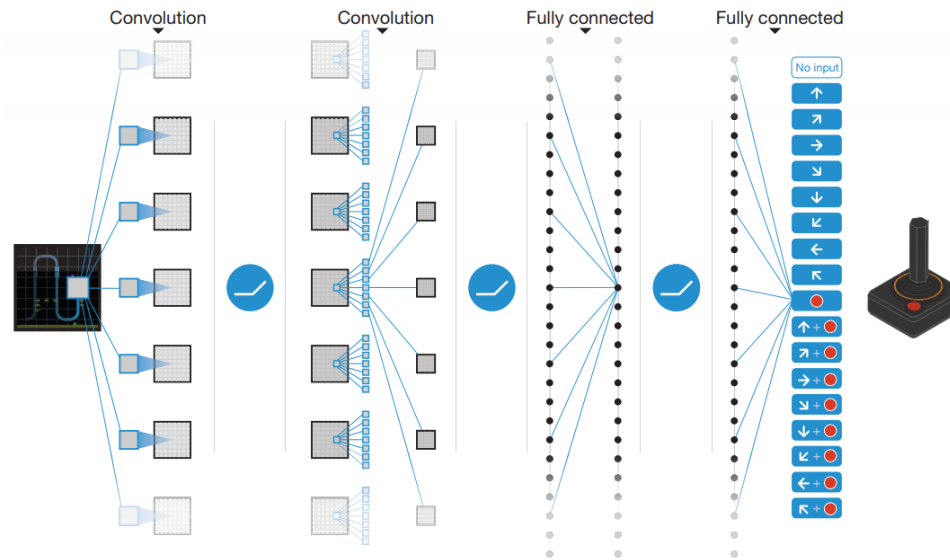
`{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com`




Deep Reinforcement Learning

■ Deep reinforcement learning

- approximate the value or policy using deep neural networks
- solving complex decision problem end-to-end



Deep Reinforcement Learning



谢 谢!

